

High-Impact Lending Decisions

September 14, 2024

In the banking sector, the decision to grant loans is a critical process that hinges on a meticulous evaluation of various customer variables. Factors such as credit score, income stability, employment history, existing debt levels, and financial behavior are thoroughly analyzed to assess the credit-worthiness of potential borrowers. This rigorous assessment is essential for minimizing default risks and ensuring the financial stability of the institution. Banks employ sophisticated statistical models and data analytics to make informed lending decisions, which helps in predicting the likelihood of repayment and identifying potential red flags.

The necessity of such a comprehensive evaluation process cannot be overstated. By effectively managing and mitigating risks, banks protect their capital reserves, thereby maintaining liquidity and operational viability. Furthermore, this approach enables banks to extend credit to reliable customers, fostering economic expansion and enhancing customer relationships. Through prudent lending practices, banks support both individual financial goals and broader economic growth. The systematic approach to loan approval, incorporating a balance of risk and opportunity, is indispensable for the sustainable growth and profitability of financial institutions. It ensures that banks can continue to play a pivotal role in the financial ecosystem by facilitating investment, consumption, and economic stability.

0.1 The Libraries Required

```
[2]: !pip install xgboost
```

```
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: xgboost in
c:\users\sayak\appdata\roaming\python\python311\site-packages (2.1.0)
Requirement already satisfied: numpy in e:\anaconda\lib\site-packages (from
xgboost) (1.24.3)
Requirement already satisfied: scipy in e:\anaconda\lib\site-packages (from
xgboost) (1.11.1)
```

```
[94]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from xgboost import XGBClassifier as xgb
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import r2_score
from scipy.stats import chi2_contingency
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, \
    precision_recall_fscore_support
import warnings
import os

```

0.2 The Dataset

In this project, I am going to work with two datasets, which contains information regarding 50k+ customers of some bank along with their personal information like, marital status, education, credit score and so on.

```

[4]: a1 = pd.read_excel("case_study1.xlsx")
     a2 = pd.read_excel("case_study2.xlsx")

```

```

[5]: df1 = a1.copy()
     df2 = a2.copy()

```

```

[6]: print(df1.shape)
     print(df2.shape)

```

```
(51336, 26)
```

```
(51336, 62)
```

```

[7]: df1.head()

```

```

[7]: PROSPECTID  Total_TL  Tot_Closed_TL  Tot_Active_TL  Total_TL_opened_L6M  \
0             1         5             4             1             0
1             2         1             0             1             0
2             3         8             0             8             1
3             4         1             0             1             1
4             5         3             2             1             0

```

```

     Tot_TL_closed_L6M  pct_tl_open_L6M  pct_tl_closed_L6M  pct_active_tl  \
0                   0             0.000             0.0             0.200
1                   0             0.000             0.0             1.000
2                   0             0.125             0.0             1.000
3                   0             1.000             0.0             1.000
4                   0             0.000             0.0             0.333

```

```

     pct_closed_tl  ...  CC_TL  Consumer_TL  Gold_TL  Home_TL  PL_TL  \
0             0.800  ...    0           0           1           0           4
1             0.000  ...    0           1           0           0           0
2             0.000  ...    0           6           1           0           0
3             0.000  ...    0           0           0           0           0
4             0.667  ...    0           0           0           0           0

```

	Secured_TL	Unsecured_TL	Other_TL	Age_Oldest_TL	Age_Newest_TL
0	1	4	0	72	18
1	0	1	0	7	7
2	2	6	0	47	2
3	0	1	1	5	5
4	3	0	2	131	32

[5 rows x 26 columns]

```
[77]: df2.head()
```

```
[77]: PROSPECTID  time_since_recent_payment  num_times_delinquent  \
0            1                549                11
1            2                47                0
2            3               302                9
4            5               583                0
5            6               245               14

max_recent_level_of_deliq  num_deliq_6mts  num_deliq_12mts  \
0                29                0                0
1                0                0                0
2               25                1                9
4                0                0                0
5              270                0                0

num_deliq_6_12mts  num_times_30p_dpd  num_times_60p_dpd  num_std  ...  \
0                0                0                0        21  ...
1                0                0                0         0  ...
2                8                0                0        10  ...
4                0                0                0        53  ...
5                0               13               11         5  ...

pct_PL_enq_L6m_of_L12m  pct_CC_enq_L6m_of_L12m  pct_PL_enq_L6m_of_ever  \
0                0.0                0.0                0.000
1                0.0                0.0                0.000
2                0.0                0.0                0.000
4                0.0                0.0                0.000
5                1.0                0.0                0.429

pct_CC_enq_L6m_of_ever  HL_Flag  GL_Flag  last_prod_enq2  first_prod_enq2  \
0                0.0          1          0              PL              PL
1                0.0          0          0  ConsumerLoan  ConsumerLoan
2                0.0          1          0  ConsumerLoan        others
4                0.0          0          0              AL              AL
5                0.0          1          0  ConsumerLoan              PL

Credit_Score  Approved_Flag
```

0	696	P2
1	685	P2
2	693	P2
4	753	P1
5	668	P3

[5 rows x 54 columns]

Clearly the datasets needed to be cleaned and pre-processed before it can be used for modelling. So, we shall start with possible “Exploratory Data Analysis”

0.3 Exploratory Data Analysis

Remove Nulls When companies gathers data, instead of keeping a cell blank or absurd, it tends to input a pre-decided value. In this scenario, the value “-99999” is one of such. So we shall try to remove or impute this value considering it as NULL/NA value.

```
[9]: #In case of first data, df1. Only one column has -99999 values that is
      ↪ "Age_Oldest_TL" column
df = df1.loc[df1['Age_Oldest_TL'] == -99999]
df.shape[0]
```

[9]: 40

Only 40 rows has this value, so we can just remove those rows as our dataset is huge this small number wont affect.

```
[10]: df1 = df1.loc[df1['Age_Oldest_TL'] != -99999]
df1.shape
```

[10]: (51296, 26)

For the second dataset, df2 we shall adopt a scheme for handling this NA value:- If the number of “-99999” values are more than 10000 we shall drop the column, else we shall drop the number of rows. At the end if we can retain more than 80% of the data, we are good to go.

```
[11]: # Here we figure out the columns that are to be removed with this scheme

columns_to_be_removed = []
for i in df2.columns:
    if df2.loc[df2[i] == -99999].shape[0] > 10000:
        columns_to_be_removed.append(i)
```

```
[12]: print(columns_to_be_removed)
df2 = df2.drop(columns_to_be_removed, axis =1)
```

```
['time_since_first_delinquency', 'time_since_recent_delinquency',
'max_delinquency_level', 'max_deliq_6mts', 'max_deliq_12mts', 'CC_utilization',
'PL_utilization', 'max_unsec_exposure_inPct']
```

```
[13]: df2.shape
```

```
[13]: (51336, 54)
```

```
[14]: # To remove the corresponding rows from the dataset
```

```
for i in df2.columns:  
    df2 = df2.loc[df2[i] != -99999]
```

```
[15]: df2.shape
```

```
[15]: (42066, 54)
```

Ultimately, more than 80% of the data is retained after removal of NULL values. So we shall proceed

Now, our next step is to merge the datasets with respect to some columns

```
[16]: # To figure out any common column in both
```

```
for i in list(df1.columns):  
    if i in list(df2.columns):  
        print(i)
```

PROSPECTID

```
[17]: # We shall merge the two datasets based on this column
```

```
# We will use inner join as we need all the common data in these two datasets.  
df = pd.merge ( df1, df2, how = 'inner', left_on = ['PROSPECTID'], right_on =  
    ↪ ['PROSPECTID'] )  
df.shape
```

```
[17]: (42064, 79)
```

```
[18]: # Number of categorical variables
```

```
for i in df.columns:  
    if df[i].dtype=='object':  
        print(i)
```

MARITALSTATUS

EDUCATION

GENDER

last_prod_enq2

first_prod_enq2

Approved_Flag

```
[19]: for i in ['MARITALSTATUS', 'EDUCATION', 'GENDER', 'last_prod_enq2',  
    ↪ 'first_prod_enq2']:  
        chi2, pval, _, _ = chi2_contingency(pd.crosstab(df[i], df['Approved_Flag']))  
        print(i, '---', pval)
```

```
MARITALSTATUS --- 3.578180861038862e-233
EDUCATION --- 2.6942265249737532e-30
GENDER --- 1.907936100186563e-05
last_prod_enq2 --- 0.0
first_prod_enq2 --- 7.84997610555419e-287
```

Since all the categorical variables has pvalue < 0.05 , we will accept all

```
[20]: # Now we shall check multicollinearity in the numerical columns
numeric_columns=[]
for i in df.columns:
    if df[i].dtype != 'object':
        numeric_columns.append(i)
```

```
[21]: vif_data=df[numeric_columns]
total_columns=vif_data.shape[1]
columns_to_be_kept=[]
column_index=0
```

```
[22]: for i in range (0,total_columns):

    vif_value = variance_inflation_factor(vif_data, column_index)

    if vif_value <= 6:
        columns_to_be_kept.append( numeric_columns[i] )
        column_index = column_index+1

    else:
        print([numeric_columns[i]])
        vif_data = vif_data.drop([ numeric_columns[i] ] , axis=1)
```

```
E:\Anaconda\Lib\site-packages\statsmodels\stats\outliers_influence.py:198:
```

```
RuntimeWarning: divide by zero encountered in scalar divide
```

```
vif = 1. / (1. - r_squared_i)
```

```
['Total_TL']
```

```
E:\Anaconda\Lib\site-packages\statsmodels\stats\outliers_influence.py:198:
```

```
RuntimeWarning: divide by zero encountered in scalar divide
```

```
vif = 1. / (1. - r_squared_i)
```

```
['Tot_Closed_TL']
```

```
['Tot_Active_TL']
```

```
['Total_TL_opened_L6M']
```

```
['Tot_TL_closed_L6M']
```

```
E:\Anaconda\Lib\site-packages\statsmodels\stats\outliers_influence.py:198:
```

```
RuntimeWarning: divide by zero encountered in scalar divide
```

```
vif = 1. / (1. - r_squared_i)
```

```

['pct_active_tl']
['pct_closed_tl']
['Total_TL_opened_L12M']
['pct_tl_open_L12M']

E:\Anaconda\Lib\site-packages\statsmodels\stats\outliers_influence.py:198:
RuntimeWarning: divide by zero encountered in scalar divide
    vif = 1. / (1. - r_squared_i)

['Auto_TL']
['Consumer_TL']
['Gold_TL']
['num_times_delinquent']

E:\Anaconda\Lib\site-packages\statsmodels\stats\outliers_influence.py:198:
RuntimeWarning: divide by zero encountered in scalar divide
    vif = 1. / (1. - r_squared_i)

['num_deliq_6mts']
['num_deliq_12mts']
['num_times_30p_dpd']
['num_std']
['num_std_6mts']
['num_dbt_6mts']
['num_lss_6mts']
['tot_enq']
['CC_enq']
['CC_enq_L6m']
['PL_enq']
['PL_enq_L6m']
['enq_L12m']
['enq_L6m']
['AGE']
['pct_of_active_TLs_ever']
['pct_opened_TLs_L6m_of_L12m']
['pct_PL_enq_L6m_of_L12m']
['pct_CC_enq_L6m_of_L12m']
['Credit_Score']

```

```

[23]: # The remaining columns are :-
      print(columns_to_be_kept)

```

```

['PROSPECTID', 'pct_tl_open_L6M', 'pct_tl_closed_L6M', 'Tot_TL_closed_L12M',
'pct_tl_closed_L12M', 'Tot_Missed_Pmnt', 'CC_TL', 'Home_TL', 'PL_TL',
'Secured_TL', 'Unsecured_TL', 'Other_TL', 'Age_Oldest_TL', 'Age_Newest_TL',
'time_since_recent_payment', 'max_recent_level_of_delinq', 'num_deliq_6_12mts',
'num_times_60p_dpd', 'num_std_12mts', 'num_sub', 'num_sub_6mts',
'num_sub_12mts', 'num_dbt', 'num_dbt_12mts', 'num_lss', 'num_lss_12mts',
'recent_level_of_delinq', 'CC_enq_L12m', 'PL_enq_L12m', 'time_since_recent_enq',
'enq_L3m', 'NETMONTHLYINCOME', 'Time_With_Curr_Empr', 'pct_currentBal_all_TL',

```

```
'CC_Flag', 'PL_Flag', 'pct_PL_enq_L6m_of_ever', 'pct_CC_enq_L6m_of_ever',
'HL_Flag', 'GL_Flag']
```

Finally we shall test for difference in means, by performing ANOVA test

```
[24]: # check Anova for columns_to_be_kept

from scipy.stats import f_oneway

columns_to_be_kept_numerical = []

for i in columns_to_be_kept:
    a = list(df[i])
    b = list(df['Approved_Flag'])

    group_P1 = [value for value, group in zip(a, b) if group == 'P1']
    group_P2 = [value for value, group in zip(a, b) if group == 'P2']
    group_P3 = [value for value, group in zip(a, b) if group == 'P3']
    group_P4 = [value for value, group in zip(a, b) if group == 'P4']

    f_statistic, p_value = f_oneway(group_P1, group_P2, group_P3, group_P4)

    if p_value <= 0.05:
        columns_to_be_kept_numerical.append(i)
```

So the final set of features are :-

```
[25]: features = columns_to_be_kept_numerical + ['MARITALSTATUS', 'EDUCATION', 'GENDER', 'last_prod_enq2', 'first_prod_enq2']
df = df[features + ['Approved_Flag']]
df.describe()
```

```
[25]:
```

	PROSPECTID	pct_tl_open_L6M	pct_tl_closed_L6M	Tot_TL_closed_L12M	\
count	42064.000000	42064.000000	42064.000000	42064.000000	
mean	25649.827477	0.179032	0.097783	0.825504	
std	14844.173396	0.278043	0.210957	1.537208	
min	1.000000	0.000000	0.000000	0.000000	
25%	12776.750000	0.000000	0.000000	0.000000	
50%	25706.500000	0.000000	0.000000	0.000000	
75%	38518.250000	0.333000	0.100000	1.000000	
max	51336.000000	1.000000	1.000000	33.000000	

	pct_tl_closed_L12M	Tot_Missed_Pmnt	CC_TL	Home_TL	\
count	42064.000000	42064.000000	42064.000000	42064.000000	
mean	0.160365	0.525746	0.145921	0.076241	
std	0.258831	1.106442	0.549314	0.358582	
min	0.000000	0.000000	0.000000	0.000000	

25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	0.250000	1.000000	0.000000	0.000000
max	1.000000	34.000000	27.000000	10.000000

	PL_TL	Secured_TL	...	time_since_recent_enq	enq_L3m \
count	42064.000000	42064.000000	...	42064.000000	42064.000000
mean	0.328000	2.921334	...	264.854507	1.230458
std	0.916368	6.379764	...	466.585002	2.069461
min	0.000000	0.000000	...	0.000000	0.000000
25%	0.000000	0.000000	...	9.000000	0.000000
50%	0.000000	1.000000	...	79.000000	1.000000
75%	0.000000	3.000000	...	302.000000	2.000000
max	29.000000	235.000000	...	4768.000000	42.000000

	NETMONTHLYINCOME	Time_With_Curr_Empr	CC_Flag	PL_Flag \
count	4.206400e+04	42064.000000	42064.000000	42064.000000
mean	2.692990e+04	110.345783	0.102962	0.193063
std	2.084300e+04	75.629967	0.303913	0.394707
min	0.000000e+00	0.000000	0.000000	0.000000
25%	1.800000e+04	61.000000	0.000000	0.000000
50%	2.400000e+04	92.000000	0.000000	0.000000
75%	3.100000e+04	131.000000	0.000000	0.000000
max	2.500000e+06	1020.000000	1.000000	1.000000

	pct_PL_enq_L6m_of_ever	pct_CC_enq_L6m_of_ever	HL_Flag \
count	42064.000000	42064.000000	42064.000000
mean	0.195497	0.064186	0.252235
std	0.367414	0.225989	0.434300
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	0.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000

	GL_Flag
count	42064.000000
mean	0.056580
std	0.231042
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 38 columns]

0.3.1 Label Encoding

Now we shall encode the categorical variables so that it will be easier to work with during modelling.

```
[26]: for i in df.columns:
        if df[i].dtype=='object':
            print(i)
```

```
MARITALSTATUS
EDUCATION
GENDER
last_prod_enq2
first_prod_enq2
Approved_Flag
```

In the Education column, the variables are ordinal, but the other columns doesn't have such ordinal characteristics

```
[27]: # We manually encoded the levels in EDUCATION column

df.loc[df['EDUCATION'] == 'SSC', ['EDUCATION']] = 1
df.loc[df['EDUCATION'] == '12TH', ['EDUCATION']] = 2
df.loc[df['EDUCATION'] == 'GRADUATE', ['EDUCATION']] = 3
df.loc[df['EDUCATION'] == 'UNDER GRADUATE', ['EDUCATION']] = 3
df.loc[df['EDUCATION'] == 'POST-GRADUATE', ['EDUCATION']] = 4
df.loc[df['EDUCATION'] == 'OTHERS', ['EDUCATION']] = 1
df.loc[df['EDUCATION'] == 'PROFESSIONAL', ['EDUCATION']] = 5
df['EDUCATION'].value_counts()
df['EDUCATION'] = df['EDUCATION'].astype(int)
```

```
[28]: df_encoded = pd.get_dummies(df, columns=['MARITALSTATUS', 'GENDER', 'last_prod_enq2', 'first_prod_enq2'])
```

```
[29]: df_encoded.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 42064 entries, 0 to 42063
Data columns (total 56 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   PROSPECTID                            42064 non-null  int64
1   pct_tl_open_L6M                       42064 non-null  float64
2   pct_tl_closed_L6M                     42064 non-null  float64
3   Tot_TL_closed_L12M                    42064 non-null  int64
4   pct_tl_closed_L12M                    42064 non-null  float64
5   Tot_Missed_Pmnt                       42064 non-null  int64
6   CC_TL                                 42064 non-null  int64
7   Home_TL                               42064 non-null  int64
8   PL_TL                                 42064 non-null  int64
9   Secured_TL                            42064 non-null  int64
```

10	Unsecured_TL	42064	non-null	int64
11	Other_TL	42064	non-null	int64
12	Age_Oldest_TL	42064	non-null	int64
13	Age_Newest_TL	42064	non-null	int64
14	time_since_recent_payment	42064	non-null	int64
15	max_recent_level_of_delinq	42064	non-null	int64
16	num_delinq_6_12mts	42064	non-null	int64
17	num_times_60p_dpd	42064	non-null	int64
18	num_std_12mts	42064	non-null	int64
19	num_sub	42064	non-null	int64
20	num_sub_6mts	42064	non-null	int64
21	num_sub_12mts	42064	non-null	int64
22	num_dbt	42064	non-null	int64
23	num_dbt_12mts	42064	non-null	int64
24	num_lss	42064	non-null	int64
25	recent_level_of_delinq	42064	non-null	int64
26	CC_enq_L12m	42064	non-null	int64
27	PL_enq_L12m	42064	non-null	int64
28	time_since_recent_enq	42064	non-null	int64
29	enq_L3m	42064	non-null	int64
30	NETMONTHLYINCOME	42064	non-null	int64
31	Time_With_Curr_Empr	42064	non-null	int64
32	CC_Flag	42064	non-null	int64
33	PL_Flag	42064	non-null	int64
34	pct_PL_enq_L6m_of_ever	42064	non-null	float64
35	pct_CC_enq_L6m_of_ever	42064	non-null	float64
36	HL_Flag	42064	non-null	int64
37	GL_Flag	42064	non-null	int64
38	EDUCATION	42064	non-null	int32
39	Approved_Flag	42064	non-null	object
40	MARITALSTATUS_Married	42064	non-null	bool
41	MARITALSTATUS_Single	42064	non-null	bool
42	GENDER_F	42064	non-null	bool
43	GENDER_M	42064	non-null	bool
44	last_prod_enq2_AL	42064	non-null	bool
45	last_prod_enq2_CC	42064	non-null	bool
46	last_prod_enq2_ConsumerLoan	42064	non-null	bool
47	last_prod_enq2_HL	42064	non-null	bool
48	last_prod_enq2_PL	42064	non-null	bool
49	last_prod_enq2_others	42064	non-null	bool
50	first_prod_enq2_AL	42064	non-null	bool
51	first_prod_enq2_CC	42064	non-null	bool
52	first_prod_enq2_ConsumerLoan	42064	non-null	bool
53	first_prod_enq2_HL	42064	non-null	bool
54	first_prod_enq2_PL	42064	non-null	bool
55	first_prod_enq2_others	42064	non-null	bool

dtypes: bool(16), float64(5), int32(1), int64(33), object(1)

memory usage: 13.3+ MB

0.3.2 Model Fitting

0.3.3 We shall consider, 3 classification algorithms, namely :-

Decision Tree

```
[30]: # Train-Test Split
from sklearn.tree import DecisionTreeClassifier
y = df_encoded['Approved_Flag']
x = df_encoded.drop ( ['Approved_Flag'], axis = 1 )
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
↳ random_state=42)
```

```
[31]: dt_model = DecisionTreeClassifier(max_depth=20, min_samples_split=10)
dt_model.fit(x_train, y_train)
y_pred = dt_model.predict(x_test)

accuracy = accuracy_score(y_test, y_pred)
print ()
print(f"Accuracy: {accuracy:.2f}")
print ()
```

Accuracy: 0.71

```
[32]: precision, recall, f1_score, _ = precision_recall_fscore_support(y_test, y_pred)

for i, v in enumerate(['p1', 'p2', 'p3', 'p4']):
    print(f"Class {v}:")
    print(f"Precision: {precision[i]}")
    print(f"Recall: {recall[i]}")
    print(f"F1 Score: {f1_score[i]}")
    print()
```

Class p1:
Precision: 0.7386934673366834
Recall: 0.7248520710059172
F1 Score: 0.7317073170731708

Class p2:
Precision: 0.8098918083462133
Recall: 0.8309217046580774
F1 Score: 0.820271989042168

Class p3:
Precision: 0.3396694214876033
Recall: 0.31018867924528304
F1 Score: 0.3242603550295859

Class p4:

Precision: 0.6366279069767442
Recall: 0.6384839650145773
F1 Score: 0.6375545851528384

Random Forest

```
[33]: rf_classifier = RandomForestClassifier(n_estimators = 200, random_state=42)
      rf_classifier.fit(x_train, y_train)
      y_pred = rf_classifier.predict(x_test)
      accuracy = accuracy_score(y_test, y_pred)
      print ()
      print(f'Accuracy: {accuracy}')
      print ()
      precision, recall, f1_score, _ = precision_recall_fscore_support(y_test, y_pred)
      for i, v in enumerate(['p1', 'p2', 'p3', 'p4']):
          print(f"Class {v}:")
          print(f"Precision: {precision[i]}")
          print(f"Recall: {recall[i]}")
          print(f"F1 Score: {f1_score[i]}")
          print()
```

Accuracy: 0.7654819921549982

Class p1:

Precision: 0.846517119244392
Recall: 0.7071005917159763
F1 Score: 0.7705534658785599

Class p2:

Precision: 0.7953968522592655
Recall: 0.931615460852329
F1 Score: 0.8581340149716997

Class p3:

Precision: 0.4423076923076923
Recall: 0.20830188679245282
F1 Score: 0.2832221652129297

Class p4:

Precision: 0.723136495643756
Recall: 0.7259475218658892
F1 Score: 0.7245392822502424

XGBoost

```
[34]: from sklearn.preprocessing import LabelEncoder
xgb_classifier = xgb(objective='multi:softmax', num_class=4)
y = df_encoded['Approved_Flag']
x = df_encoded.drop ( ['Approved_Flag'], axis = 1 )
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)
x_train, x_test, y_train, y_test = train_test_split(x, y_encoded, test_size=0.
↪2, random_state=42)
xgb_classifier.fit(x_train, y_train)
y_pred = xgb_classifier.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)
print ()
print(f'Accuracy: {accuracy:.2f}')
print ()
precision, recall, f1_score, _ = precision_recall_fscore_support(y_test, y_pred)
for i, v in enumerate(['p1', 'p2', 'p3', 'p4']):
    print(f"Class {v}:")
    print(f"Precision: {precision[i]}")
    print(f"Recall: {recall[i]}")
    print(f"F1 Score: {f1_score[i]}")
    print()
```

Accuracy: 0.78

Class p1:

Precision: 0.8238993710691824

Recall: 0.7751479289940828

F1 Score: 0.798780487804878

Class p2:

Precision: 0.826157158234661

Recall: 0.9127849355797819

F1 Score: 0.8673133063376967

Class p3:

Precision: 0.4634433962264151

Recall: 0.29660377358490564

F1 Score: 0.3617119190059825

Class p4:

Precision: 0.7270973963355835

Recall: 0.7327502429543246

F1 Score: 0.7299128751210068

Here, we observe that the precision for p3 in all cases is very low. To fix this problem we shall try to do some Hyperparameter tuning.

```
[36]: param_grid={
        'colsample_bytree':[0.1,0.3,0.5,0.7,0.9],
        'learning_rate':[0.001,0.01,0.1,1],
        'max_depth':[3,5,8,10],
        'alpha':[1,10,100],
        'n_estimators':[10,50,100]
    }

    index = 0

    answers_grid = {
        'combination'      : [],
        'train_Accuracy'    : [],
        'test_Accuracy'     : [],
        'colsample_bytree'  : [],
        'learning_rate'     : [],
        'max_depth'         : [],
        'alpha'             : [],
        'n_estimators'      : []
    }
```

```
[64]: y = df_encoded['Approved_Flag']
      x = df_encoded.drop ( ['Approved_Flag'], axis = 1 )
```

```
[57]: rs=GridSearchCV(xgb_classifier,param_grid=param_grid,cv=5,n_jobs=-1,verbose=4)
```

```
[58]: model=rs.fit(x,y)
```

Fitting 5 folds for each of 720 candidates, totalling 3600 fits

```
[43]: from sklearn import preprocessing
      enc=preprocessing.LabelEncoder()
```

```
[44]: y=enc.fit_transform(y)
```

```
[73]: model.best_score_
```

```
[73]: 0.7757700293892245
```

```
[76]: xgb_classifier = xgb(objective='multi:softmax',
                          num_class=4,
                          alpha=10,
                          colsample_bytree=0.7,
                          learning_rate=0.1,
                          max_depth=8,
                          n_estimators=50
                          )
```

```

y = df_encoded['Approved_Flag']
x = df_encoded.drop ( ['Approved_Flag'], axis = 1 )
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)
x_train, x_test, y_train, y_test = train_test_split(x, y_encoded, test_size=0.
    ↪2, random_state=42)
xgb_classifier.fit(x_train, y_train)
y_pred = xgb_classifier.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)
print ()
print(f'Accuracy: {accuracy:.2f}')
print ()
precision, recall, f1_score, _ = precision_recall_fscore_support(y_test, y_pred)
for i, v in enumerate(['p1', 'p2', 'p3', 'p4']):
    print(f"Class {v}:")
    print(f"Precision: {precision[i]}")
    print(f"Recall: {recall[i]}")
    print(f"F1 Score: {f1_score[i]}")
    print()

```

Accuracy: 0.78

Class p1:
Precision: 0.8448660714285714
Recall: 0.7465483234714004
F1 Score: 0.7926701570680629

Class p2:
Precision: 0.8080234159779615
Recall: 0.9302279484638256
F1 Score: 0.8648300009214044

Class p3:
Precision: 0.4658753709198813
Recall: 0.2369811320754717
F1 Score: 0.3141570785392696

Class p4:
Precision: 0.7352657004830918
Recall: 0.7395529640427599
F1 Score: 0.7374031007751939

For P3 we shall bootstrap samples from the data as this consistent low precision from all model might be a symptom of less representation.

[65]:


```
[65]: Approved_Flag
      P2      25452
      P3      6440
      P4      5264
      P1      4908
      Name: count, dtype: int64
```

```
[89]: from sklearn.multiclass import OneVsRestClassifier

      ovr = OneVsRestClassifier(xgb_classifier)
      y = df_encoded['Approved_Flag']
      x = df_encoded.drop ( ['Approved_Flag'], axis = 1 )
      label_encoder = LabelEncoder()
      y_encoded = label_encoder.fit_transform(y)
      x_train, x_test, y_train, y_test = train_test_split(x, y_encoded, test_size=0.
      ↪2, random_state=42)
      ovr.fit(x_train, y_train)
      y_pred=ovr.predict(x_test)
```

```
[91]: accuracy = accuracy_score(y_test, y_pred)
      print ()
      print(f'Accuracy: {accuracy:.2f}')
      print ()
      precision, recall, f1_score, _ = precision_recall_fscore_support(y_test, y_pred)
      for i, v in enumerate(['p1', 'p2', 'p3', 'p4']):
          print(f"Class {v}:")
          print(f"Precision: {precision[i]}")
          print(f"Recall: {recall[i]}")
          print(f"F1 Score: {f1_score[i]}")
          print()
```

Accuracy: 0.78

Class p1:
Precision: 0.844789356984479
Recall: 0.7514792899408284
F1 Score: 0.7954070981210856

Class p2:
Precision: 0.8059548254620124
Recall: 0.933597621407334
F1 Score: 0.8650932133345579

Class p3:
Precision: 0.47572815533980584
Recall: 0.2218867924528302
F1 Score: 0.30262480699948535

```
Class p4:  
Precision: 0.7330791229742613  
Recall: 0.7473275024295433  
F1 Score: 0.7401347449470644
```

0.4 Conclusion

We are getting an average accuracy of 78% on test data which seem pretty good but, the classification accuracy for p3 is poor. To tackle this issue, we performed hyperparameter tuning and even bootstrapped data to balanced the imbalanced situation but the problem prevails. In the upcoming semester we might continue this project with the help of clustering algorithms and update the file.

[]: